


Qwen3-1.7B Math C/D0 — Model1 and Model2 dynamics

C updates both submodels; D0 intentionally keeps Model 1 unchanged.

C: weak-logit mixture, both trainable



D0: matched-scale no-weak, Model 1 frozen

